

Engineering Management

Open-Sourcing the Institution

~~Working Paper — Draft~~

31/01/2026

January 31, 2026

Summary

Organizations rely on opacity — implicit systems of credit, funding, access, and decision-making understood by insiders but invisible to outsiders and unaccountable to anyone. Machine intelligence is ending this arrangement. AI systems already reconstruct what institutions are actually doing regardless of what those institutions claim. This paper argues that organizations face a choice: adopt structured transparency on their own terms, or have it imposed from outside. We describe an architecture for *open-sourcing the institution itself* — governance, resource flows, decision history, and credit assignment — using version-controlled repositories, automated integrity checks, and machine-readable organizational state. Management becomes engineering: both positive reinforcement (recognition, credit, visibility) and negative reinforcement (integrity checks, conflict detection, staleness warnings) are code in the repository, auditable by anyone. We present the design of a research lab built on these principles and discuss its applicability to any organization that manages resources and produces work products.

Revision History

Revision	Date	Author(s)	Description
v0.1	31/01/2026	—	Initial working draft
v0.2	31/01/2026	—	Added Related Work, Limitations, Conclusion; trimmed recognition section; addressed reviewer feedback

1 Contents

2	1 Introduction	3
3	2 Related Work	3
4	3 The End of Organizational Opacity	5
5	4 The Cost of Opaque Institutions	5
6	5 Open-Sourcing the Institution	6
7	6 Architecture	7
8	6.1 The Organization Repo	7
9	6.2 Integrity Checks	8
10	6.3 Automated Recognition	9
11	6.4 The Project Management Repo	9
12	7 Relationship to Existing Institutions	10
13	8 Limitations and Open Questions	10
14	9 Conclusion	11

get rid of line numbers

15 1 Introduction

16 We need a research lab where the lab itself is open source. Not just the research outputs — the
 17 organizational structure, decision-making processes, resource allocation, and governance are
 18 all visible, forkable, and improvable by anyone. Forkability is central: if someone disagrees with
 19 how the organization operates, they can copy the entire institutional structure — governance
 20 rules, resource tracking templates, integrity checks, project management configuration — and
 21 run it differently. Organizational design becomes subject to the same evolutionary pressure as
 22 open-source software.

23 Management becomes engineering. Not because management gets reduced to well-defined
 24 correctness criteria — engineering doesn't have those either. Engineering is about tradeoffs,
 25 and so is management. The difference is that engineering tradeoffs are made explicit: docu-
 26 mented, versioned, debatable, auditable. Management tradeoffs currently live in the discretion
 27 of whoever holds a management title, invisible to the people affected by them. Making them
 28 explicit infrastructure — resource allocation rules, credit assignment criteria, accountability
 29 mechanisms as code in a repo — doesn't eliminate the tradeoffs. It makes them visible so
 30 the people affected can participate in deciding them. Both positive reinforcement (recog-
 31 nition, credit, visibility) and negative reinforcement (integrity checks, conflict detection, staleness
 32 warnings) become auditable systems rather than unaccountable discretion.

33 The lab is managed through changes to a public management project on GitHub. The rest of
 34 this paper is organized as follows. Section 2 positions this work relative to prior approaches to
 35 transparent governance. Section 3 argues that organizational opacity is becoming unsustain-
 36 able. Section 4 quantifies the economic cost of opaque institutions. Section 5 describes what
 37 it means to open-source an institution. Section 6 presents the concrete architecture. Section 7
 38 addresses the relationship to existing institutions. Section 8 discusses limitations. Section 9
 39 concludes.

40 2 Related Work

move to end

41 The idea of making organizational governance explicit and transparent has a substantial his-
 42 tory. This section positions the present work relative to prior approaches and identifies what is
 43 genuinely new.

44 **Commons governance.** Ostrom [1] established eight design principles for self-governing
 45 institutions managing shared resources: clearly defined boundaries, proportional costs and
 46 benefits, collective choice arrangements, monitoring, graduated sanctions, conflict resolution
 47 mechanisms, minimal recognition of rights to organize, and nested enterprises. The archi-
 48 tecture described in this paper instantiates several of these principles — integrity checks as
 49 monitoring, governance-as-code as collective choice arrangements, the recursive tech tree as
 50 nested enterprises — in a version-controlled, machine-readable form. Ostrom's framework
 51 provides the theoretical grounding; we provide a concrete implementation using modern soft-
 52 ware infrastructure.

53 **Self-management frameworks.** Holacracy [2] encodes governance rules explicitly, replac-
 54 ing traditional management hierarchy with a constitution that defines roles, accountabilities,

and decision-making processes. Laloux [3] surveys organizations operating with radical transparency, including Buurtzorg and Morning Star. Our approach shares the goal of making governance rules explicit but differs in two ways: the rules are stored in Git (not a proprietary constitution), making them forkable and diffable; and machine intelligence is used to check organizational state against the rules, rather than relying solely on human process adherence. Holacracy’s adoption difficulties — Zappos being the most prominent case — also illustrate that rigid governance frameworks can fail when they don’t accommodate organizational ambiguity. Our approach is deliberately flexible: governance documents are prose, not a fixed schema.

Decentralized Autonomous Organizations (DAOs). DAOs encode governance in smart contracts on a blockchain, making institutional rules machine-readable and forkable [4]. The goals overlap significantly with ours. The key differences are infrastructural: our approach uses Git, which is universally understood by developers, does not require tokens or chain infrastructure, and integrates with existing development tools. DAOs also have a well-documented history of governance failures — including the 2016 DAO exploit, where transparent machine-readable rules were used against the organization, and persistent voter apathy in token-based governance. We take these as evidence that governance transparency is necessary but not sufficient; the integrity checks and recognition system described in Section 6 are designed to address some of these failure modes.

Open-source governance. The open-source software movement provides decades of evidence on transparent organizational structures [5]. Research on open-source governance shows both the power of open contribution models and their failure modes: hidden hierarchies, burnout among maintainers, and difficulty with accountability when no one is formally in charge. Buffer’s radical transparency experiment (public salaries and open financials since 2013) provides real-world data on organizational transparency in a corporate setting. Valve’s nominally flat structure produced documented problems with clique-based resource allocation despite apparent openness — a cautionary tale that motivates our integrity checks, which are designed precisely to surface hidden power structures that persist even under nominal transparency.

The transparency paradox. Bernstein [6] found that factory workers were more productive when given privacy from management observation, showing that transparency can reduce performance by inducing hiding behavior. This is an important counterpoint, but the mechanism matters: Bernstein’s workers had no control over the transparency system and no ability to use it themselves. The transparency was asymmetric — management watched workers, not the reverse. This project is specifically an attempt at solving that problem by making transparency symmetric: machine intelligence gives every member the same ability to read and audit organizational state that was previously reserved for management. The remaining risk — that even symmetric transparency produces chilling effects on exploratory work — is addressed in Section 8.

Algorithmic management. Kellogg et al. [7] provide a systematic review of algorithmic management in the workplace, documenting how automated systems can reproduce or amplify existing power asymmetries. The automated recognition system described in Section 6.3 is an instance of algorithmic management and inherits these risks. We address this by making the recognition infrastructure itself open and auditable — the criteria are code, not black-box algorithms — but acknowledge that this does not eliminate all risks (see Section 8).

100 3 The End of Organizational Opacity

101 Machine intelligence is making organizational opacity increasingly difficult to maintain. AI sys-
102 tems can aggregate public filings, leaked documents, employment records, financial data, and
103 published research into coherent pictures of what organizations are actually doing — regard-
104 less of what those organizations say they're doing. The trend is toward increasing institutional
105 legibility: not that all institutions are already fully transparent to machines, but that the cost of
106 piercing opacity is dropping rapidly for anyone motivated to do so.

107 Organizations that don't open up voluntarily will have their opacity pierced by others. Compa-
108 nies are deploying blockchain-based supply chain verification because they can't verify what
109 opaque suppliers claim about sourcing and labor practices — a market projected to grow
110 from \$2.9 billion to \$44.3 billion by 2034 [8]. The era of "trust us, it works" is ending across
111 investment — venture investors now conduct spot-checks by contacting claimed customers di-
112 rectly and demanding bank statements behind reported revenue, because they can no longer
113 take institutional self-reporting at face value [9]. Self-insured employers are building internal
114 claims analytics because they can't trust insurer-reported numbers — medical billing errors ac-
115 count for an estimated \$68 billion in unnecessary healthcare costs annually, and half of those
116 who dispute denied claims get them reversed, suggesting the denials were wrong in the first
117 place [10].

118 This is parallel construction at organizational scale — building an independent picture of an
119 institution's internals from external data, analogous to the law enforcement technique of re-
120 constructing evidence from public sources to avoid revealing classified intelligence methods.
121 When institutions are opaque, anyone with resources builds their own apparatus to see through
122 them. But it's not only the powerful who gain this capability. The same AI tools that let investors
123 reconstruct an institution's internals are available to employees, researchers, and the public.
124 An individual can now aggregate an organization's public filings, job postings, Glassdoor re-
125 views, published outputs, and financial records into a picture that used to require a dedicated
126 analyst team. The opacity that once protected mismanagement from scrutiny above also pro-
127 tected it from scrutiny below. Both protections are ending at the same time.

128 The result is a world where the people running opaque organizations are the last to know what
129 everyone else — funders, competitors, and their own staff — already sees. The question is
130 not whether your organization will become transparent. It is whether you will be the one who
131 decides how.

132 4 The Cost of Opaque Institutions

133 The economic costs of organizational opacity are substantial, though precise measurement is
134 difficult. The estimates below draw from a mix of peer-reviewed research, industry surveys,
135 and trade publications of varying methodological rigor. Table 2 summarizes the sources and
136 their limitations.

137 By one estimate, excess management and bureaucracy cost the U.S. economy more than
138 \$3 trillion per year — roughly 17% of GDP [11]. The Fortune 500 alone may account for
139 \$480 billion in back-office inefficiencies [12]. The average firm spends between 1.3% and

140 3.3% of its total wage bill on regulatory compliance, with aggregate compliance costs ranging
 141 from \$103 billion to \$289 billion annually [13]. In one Australian study, non-administrative staff
 142 reported spending 6.4 hours per week — 16% of a standard work week — complying with
 143 internal regulation, and the cost of self-imposed internal rules exceeded the cost of government
 144 regulation by more than two to one [14].

145 Nepotism and favoritism compound these costs. Nearly 75% of employees report having
 146 worked in a toxic workplace, with poor leadership — including favoritism and lack of account-
 147 ability — cited as the top cause by 79% of them [15]. Toxic culture is 10 times more important
 148 than compensation in predicting turnover; by one estimate (originally from SHRM, cited in [16]),
 149 toxic culture cost U.S. employers nearly \$50 billion per year before the Great Resignation.
 150 When unqualified hires hold leadership roles through connections rather than competence,
 151 the remaining staff pick up the slack, leading to burnout and further turnover. The organiza-
 152 tional knowledge needed to actually run the institution is never written down — it lives in the
 153 heads of whoever happens to be politically connected enough to stay.

154 The same patterns appear in public institutions. Researchers receiving federal grants spend
 155 over 44% of their time on administrative tasks rather than research — a figure that has re-
 156 mained stubbornly consistent across FDP surveys in 2005, 2012, and 2018 [17]. Universities
 157 absorb \$6.8 billion per year in unreimbursed overhead costs [18]. Nepotism in public institu-
 158 tions creates wasteful overstaffing where unneeded positions are created to employ relatives
 159 — documented across municipalities as a drag on economic development [19]. Perceived
 160 nepotism is negatively associated with educational investment across countries, as measured
 161 by PISA scores [20]. Funding decisions happen behind closed doors. Credit assignment is
 162 political. Access to resources depends on who you know, not what you contribute.

Table 2: Summary of cost estimates cited in this section. Estimates vary in methodology,
 geography, and recency. They are presented to indicate the scale of the problem, not as
 precise measurements.

Estimate	Source	Type	Notes
\$3T excess mgmt.	Hamel (2016)	HBR opinion	Extrapolation from exemplar cos.
\$480B back-office	SSO Network	Trade pub.	Fortune 500 only
\$103–289B compliance	NBER/Cato	Working paper	U.S. firms, wage-bill method
6.4 hrs/wk internal rules	Deloitte AU	Consulting	Australian data
\$50B toxic culture	SHRM (2019)	Industry survey	Via Sull et al.
44% admin time (grants)	FDP (2018)	Survey	n = 13,000+ U.S. PIs
\$6.8B overhead	AAU	Advocacy org.	U.S. universities

163 5 Open-Sourcing the Institution

164 Most “open science” initiatives open-source the outputs (papers, data, code) but keep the
 165 institution opaque. This lab open-sources the institution itself: governance, resource flows,
 166 decision history, credit assignment. The structure is the product.

167 This isn’t all-or-nothing. An organization doesn’t have to go from political hell to fully open
 168 overnight. The move can be done gradually. Start by making governance documents public.

169 Then resource tracking. Then decision rationale. Each step is independently useful and each
 170 step makes the next one easier. The tools described here support any point on that spectrum.

171 6 Architecture

172 Here is what it looks like when an organization has no secrets from itself. Higher-level units
 173 (the lab itself, divisions within it) have two repos: a project management repo (managed by pm)
 174 and an organization repo. Individual research projects below them may only have a pm repo
 175 and their code repo. The two-repo structure is for organizational layers that need to track more
 176 than just a tech tree.

177 6.1 The Organization Repo

178 A GitHub repo that is the public record of the organization. This is where things live that don't
 179 belong in any individual project but need to be tracked: documents, small scripts, spread-
 180 sheets, test programs, analyses, meeting notes, governance decisions, and anything else that
 181 should be visible and version-controlled but doesn't need wide release as a standalone project.

182 It also includes pointers to all the other repos in the organization — the pm repo, the individual
 183 project repos, external dependencies — in docs/ so it can act as a central dispatch.

```
184 org-repo/
185 |-- docs/                                # golden copy of project state
186 |   |-- governance/                      # decision-making processes,
187 |   | roles
188 |   |   |-- research-agenda/            # what the lab works on and why
189 |   |   |-- resources/                  # grant tracking, allocation
190 |   |   records
191 |   |   '-- onboarding/                # how to join, how things work
192 |-- members/
193 |   |-- alice/                         # Alice's working directory
194 |   |   |-- notes/                     # her meeting notes, scratch work
195 |   |   |-- scripts/                  # one-off analysis scripts
196 |   |   '-- proposals/              # drafts she's working on
197 |   |-- bob/
198 |   |   |-- notes/                   # small datasets, spreadsheets
199 |   |   |-- data/                    # test programs, prototypes
200 |   |   '-- experiments/           # integrity checks (see below)
201 |   '-- ...
202 '-- checks/                            # completed work, historical
203 '-- archive/
204   records
```

205 **Member directories.** Every member gets their own directory — their workspace within the
 206 org. They can put anything there: notes, scripts, Excel sheets, draft proposals, experimental
 207 code. Their work is visible to the rest of the organization without requiring them to publish it
 208 anywhere else.

209 When a member wants to promote something to the project-wide golden copy, they open a PR
 210 moving or copying it into docs/. The PR is a request for the organization to recognize and

211 adopt the work. Other members review it, and the merge is the organization saying “yes, this
212 is now part of our shared state.”

213 **The docs/ directory.** The golden copy — the authoritative state of the project. Nothing gets
214 into docs/ without a PR. The PR history in docs/ is the lab’s institutional memory.

215 **The checks/ directory.** Automated integrity checks that project managers create and main-
216 tain. These run against the repo (via CI or on-demand) and surface issues that humans should
217 look at.

218 6.2 Integrity Checks

219 Integrity checks are to organizational state what tests are to code. When AI agents produce
220 code, tests check that the code does what it claims. When an organization produces doc-
221 uments, decisions, and resource allocations, integrity checks do the same thing: they verify
222 organizational state against the organization’s own stated rules and surface inconsistencies
223 for human judgment.

224 This applies to any organization that manages resources — public or private, academic or
225 corporate. The checks maintain a history of both automated and human-assisted audits. Every
226 check run is logged. When a human reviews an issue surfaced by a check and makes a
227 judgment call, that judgment is recorded too. Over time this builds an audit trail showing not just
228 the current state but how the organization has responded to issues as they were discovered
229 — an accountability mechanism that persists regardless of personnel changes.

230 Examples:

- 231 • **Consistency:** does the resource allocation in docs/resources/ add up? Do budget num-
232 bers match grant amounts recorded in governance decisions?
- 233 • **Staleness:** are there proposals sitting for months with no PR to docs/? Governance docu-
234 ments referencing people who are no longer active?
- 235 • **Completeness:** does every active grant have an allocation record? Does every research
236 project in the agenda have a corresponding pm instance?
- 237 • **Conflict:** contradictory statements across governance documents? Two proposals allocat-
238 ing the same funds differently?
- 239 • **Attribution:** does every work product have clear attribution? Documents referencing work
240 without crediting the contributor?

241 These checks can be simple scripts, LLM-powered analyses, or structured validators. They’re
242 not gates — they surface issues for human judgment, same as a test suite surfaces failures for
243 a developer to evaluate. The checks themselves are in the repo, subject to the same review
244 process as everything else.

245 6.3 Automated Recognition

246 The same infrastructure that detects problems also recognizes achievements. In traditional
 247 organizations, recognition flows through managers, shaped by who is visible, who is in the
 248 right social circle, who shares the manager's background. Automated recognition replaces this
 249 with a read operation on organizational state that's already being maintained.

250 Examples: milestone completion across a tech tree; a new member's first PR to docs/; sus-
 251 tained contribution over time; a PR that unblocks several downstream projects; a grant's deliv-
 252 erables all completed with the full funding-to-output chain visible.

253 Because everything in the org repo and pm repo is readable text and structured data, anyone
 254 can ask an LLM to assess contributions holistically — reading a member's directory, their PRs,
 255 the projects they've touched, the discussions they've participated in. This is closer to what a
 256 thoughtful colleague would say if they'd been paying attention to everything, which no human
 257 can do at organizational scale.

258 This avoids the problems of fixed metrics. Lines of code, commit frequency, and citation counts
 259 reward the metric instead of the work. LLM-based assessment introduces its own risks —
 260 LLMs have predictable biases toward verbose, confident, and recently-active contributors, and
 261 their evaluation can be anticipated and gamed. The game doesn't disappear; it becomes
 262 a meta-game. But because the recognition criteria, prompts, and outputs are all in the repo,
 263 anyone can see how the game is being played. The meta-game itself is auditable — if someone
 264 is optimizing for the evaluator rather than doing useful work, that pattern is visible in the same
 265 data the evaluator reads. You can write integrity checks against the meta-game the same way
 266 you write them against the organization's primary state.

267 The point is that the organization notices what its members accomplish without requiring self-
 268 promotion or managerial attention. The system does the noticing. People decide what to do
 269 with it. This is not a tool for optimizing productivity — it will surface what looks like inefficiency
 270 but is actually useful exploration, and the organization should understand that distinction.

271 6.4 The Project Management Repo

272 Separate from the org repo. Managed by `pm`. Contains `project.yaml`, `plans/`, and the tech
 273 tree. This is where work planning happens: what needs to be built, in what order, who's
 274 working on what. *what is pm? We need to link to it. We also need a demo for setting up an org*

275 The pm repo integrates with a recursive tech tree system. The lab's top-level tree connects
 276 the organizational layer to all the research projects below it. In *prescriptive* mode, a parent
 277 tree suggests work to child projects (the child opts in and can accept or decline). In *descriptive*
 278 mode, a parent tree observes child projects without directing them — a read-only aggregation
 279 that requires no coordination. *with the goal at that link.*

280 The org repo and pm repo are separate concerns: the org repo is the record of what the
 281 organization is and has done; the pm repo is the plan for what it's doing next. The pm repo
 282 will expose what looks like inefficiency — blocked work, stalled plans, idle projects. Some of
 283 that is real inefficiency. Some of it is useful exploration. The visibility lets the organization have
 284 honest conversations about what's happening.

285 7 Relationship to Existing Institutions

286 The lab is additive, not exclusionary. Researchers can participate while holding positions at
287 existing institutions. Grants can come from traditional funders and flow through traditional fiscal
288 sponsors. The transparency is about what happens with the resources, not about requiring
289 new channels.

290 As long as existing power structures persist, people will interface with governments through ex-
291 isting institutions for taxes, employment law, visa sponsorship, and compliance. A researcher
292 receiving a stipend still files taxes through their employer. A grant still goes through a fiscal
293 sponsor. The open layer sits on top of these structures, making decisions and flows visible
294 without pretending the underlying structures don't exist.

295 This also matters for funders. All else equal, an organization where a machine intelligence can
296 read and verify what's actually happening is less risky to fund than an opaque one. As this
297 kind of transparency becomes feasible, it will increasingly be expected — not because anyone
298 mandates it, but because the alternative becomes an obviously worse bet.

299 8 Limitations and Open Questions

300 **The transparency paradox.** Bernstein [6] demonstrated that factory workers were more pro-
301 ductive when given privacy from management observation. This is a genuine finding, but the
302 mechanism matters: Bernstein's workers had no control over the observation system and no
303 ability to use it themselves. Transparency was asymmetric — management watched work-
304 ers, not the reverse. In that configuration, transparency is surveillance, and surveillance in-
305 duces hiding behavior. This project is specifically an attempt at solving the problem Bernstein
306 identified, by making transparency symmetric. When every member has the same machine-
307 intelligence-assisted access to organizational state that executives and project managers have,
308 the power asymmetry that drives hiding behavior is substantially reduced. A tech worker with
309 AI tools can read governance documents, audit resource flows, and assess organizational
310 health with nearly the same capability as the people nominally in charge. The mechanism
311 described here is more symmetric than prior transparency systems precisely because ma-
312 chine intelligence makes organizational legibility available to everyone, not just to those with
313 management titles. Nevertheless, the risk of chilling effects on exploratory work remains real.
314 Members might avoid speculative projects that could look unproductive. The member direc-
315 tories are designed to mitigate this — they are personal workspaces where incomplete and
316 exploratory work is expected — but the tension between organizational legibility and individual
317 creative freedom requires ongoing attention.

318 **Legitimate uses of opacity.** Not all organizational opacity is pathological. Trade secrets,
319 negotiating positions, personnel matters, and shielding individuals from external pressure are
320 cases where some degree of opacity serves legitimate purposes. The architecture described
321 here does not require total transparency — the org can choose what is world-readable and
322 what is member-only. But the paper's argument is strongest for governance, resource alloca-
323 tion, and credit assignment, where opacity more often enables abuse than serves legitimate
324 ends.

325 **LLM gaming and bias.** The automated recognition system inherits the risks of algorithmic
326 management documented by Kellogg et al. [7]. LLMs have predictable biases and their as-
327 sessments can be gamed by those who understand the model’s preferences. Making the
328 recognition infrastructure open and auditable means the meta-game — gaming the evaluator
329 rather than doing the work — is itself visible and auditable. This does not eliminate strategic
330 behavior, but it makes each layer of strategy legible to anyone who cares to look. Formal
331 analysis of these recursive gaming surfaces is future work.

332 **No empirical validation.** This paper describes an architecture, not empirical results. The
333 claims about cost reduction, improved fairness, and organizational health are predictions, not
334 measurements. A pilot deployment (even at small scale: 5–10 people for 3–6 months) measur-
335 ing administrative overhead, perceived fairness, and contribution distribution would substan-
336 tially strengthen the argument. Formal modeling — agent-based simulation of organizational
337 dynamics under transparent vs. opaque regimes, game-theoretic analysis of strategic behav-
338 ior, principal-agent models adapted to symmetric transparency — is also future work. This is
339 a design proposal; the empirical and theoretical validation comes next.

340 **Cost figures.** The economic estimates in Section 4 draw from sources of varying rigor (see
341 Table 2). They indicate the scale of the problem but should not be treated as precise mea-
342 surements. In particular, the \$3 trillion figure from Hamel [11] is based on extrapolation from
343 exemplar companies, not a comprehensive economic study.

344 9 Conclusion

345 Organizational opacity is becoming unsustainable. Machine intelligence is making institutional
346 internals legible to outsiders whether institutions cooperate or not. The costs of the opaque in-
347 stitutions we have — excess management, bureaucratic overhead, nepotism, lost productivity
348 — are measured in trillions.

349 This paper proposes an alternative: open-source the institution itself. Not just the research
350 outputs, but the governance, resource flows, decision history, and credit assignment. The
351 concrete architecture — an organization repo with integrity checks, an automated recognition
352 system, and a recursive project management layer — is implementable with existing tools. The
353 key insight is that integrity checks are to organizational state what tests are to code: they hold
354 the institution accountable to its own stated rules.

355 Management becomes engineering. The structure is the product. The first organization to fully
356 open-source itself will make every opaque institution around it look like a liability.

References

- 358 [1] E. Ostrom, *Governing the Commons: The Evolution of Institutions for Collective Action*.
359 Cambridge University Press, 1990.
- 360 [2] B. J. Robertson, *Holacracy: The New Management System for a Rapidly Changing World*.
361 Henry Holt, 2015.
- 362 [3] F. Laloux, *Reinventing Organizations*. Nelson Parker, 2014.
- 363 [4] D. Rozas, A. Tenorio-Fornés, S. Djunadi, and J. Hassan, “Analysis of the Potentials of
364 Blockchain Technology for Commons Governance,” 2021.
- 365 [5] E. S. Raymond, *The Cathedral and the Bazaar*. O'Reilly, 1999.
- 366 [6] E. S. Bernstein, “The Transparency Paradox: A Role for Privacy in Organizational Learn-
367 ing and Operational Control,” *Administrative Science Quarterly*, vol. 57, no. 2, pp. 181–
368 216, 2012.
- 369 [7] K. C. Kellogg, M. A. Valentine, and A. Christin, “Algorithms at Work: The New Contested
370 Terrain of Control,” *Academy of Management Annals*, vol. 14, no. 1, pp. 366–410, 2020.
- 371 [8] Market.us, “Blockchain for Supply Chain Traceability Market Size,” 2024. <https://market.us/report/blockchain-for-supply-chain-traceability-market/>
- 373 [9] Triunity Capital Ventures, “Due Diligence in 2025: What Investors Really Want
374 to Know About Your AI Startup.” <https://triunitycapitalventures.com/due-diligence-in-2025-what-investors-really-want-to-know-about-your-ai-startup/>
- 376 [10] AJMC, “Survey Exposes Pervasive Billing Errors, Aggressive Tac-
377 tics in US Health Insurance,” 2024. <https://www.ajmc.com/view/survey-exposes-pervasive-billing-errors-aggressive-tactics-in-us-health-insurance>
- 379 [11] G. Hamel, “Excess Management Is Costing the U.S. \$3 Trillion Per
380 Year,” *Harvard Business Review*, 2016. <https://hbr.org/2016/09/excess-management-is-costing-the-us-3-trillion-per-year>
- 382 [12] SSO Network, “Fortune 500: \$480 Billion in Back Office Inefficiencies a Year.”
383 <https://www.ssonetwork.com/continuous-improvement-process-improvement/articles/480-billion-in-back-office-inefficiencies-how-to>
- 385 [13] NBER / Cato, “The Cost of Regulatory Compliance in the United
386 States.” <https://www.cato.org/research-briefs-economic-policy/cost-regulatory-compliance-united-states>
- 388 [14] Deloitte Australia, “Get Out of Your Own Way — Unleashing Productiv-
389 ity,” cited in *Consultancy.uk*. <https://www.consultancy.uk/news/973/deloitte-compliance-costs-australian-economy-250-billion>
- 391 [15] iHire, “Toxic Workplace Trends Report,” February 2025 ($n =$
392 2,285 across 57 industries). <https://www.ihire.com/about/press/ihire-toxic-workplace-trends-report-pr>

- 394 [16] D. Sull, C. Sull, W. Cipolli, and C. Brighenti, "Why Every Leader Needs to Worry About
395 Toxic Culture," *MIT Sloan Management Review*, March 2022. \$50B figure originally from
396 SHRM, "The High Cost of a Toxic Workplace Culture," 2019. <https://sloanreview.mit.edu/article/why-every-leader-needs-to-worry-about-toxic-culture/>
- 397
- 398 [17] S. L. Schneider, "2018 FDP Faculty Workload Survey: Research Report: Primary
399 Findings," Federal Demonstration Partnership Foundation, 2020. <https://thefdp.org/wp-content/uploads/FDP-FWS-2018-Primary-Report.pdf>
- 400
- 401 [18] Association of American Universities, "Frequently Asked Questions about Facilities and
402 Administrative Costs of Federally Sponsored University Research." <https://www.aau.edu/key-issues/frequently-asked-questions-about-facilities-and-administrative-costs>
- 403
- 404
- 405 [19] "Nepotism, Political Competition and Overemployment," *Journal of Public Economic Theory*, 2020. <https://www.tandfonline.com/doi/full/10.1080/2474736X.2020.1781542>
- 406
- 407
- 408 [20] "Nepotism, Human Capital and Economic Development," *Journal of Economic Behavior & Organization*, 2020. <https://www.sciencedirect.com/science/article/abs/pii/S0167268120304431>
- 409
- 410